

Climate Data Analysis

Task:-

- Undertake a comprehensive climate data analysis project to explore and understand historical climate patterns and trends. The objective is to derive valuable insights from climate data, enabling a better understanding of weather conditions over time.

Importing necessary libraries

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import seaborn as sns
6 import datetime
7 import warnings
8 warnings.filterwarnings('ignore')
9 sns.set()
```

Loading and Exploring the Dataset

In [2]:

```
1 df = pd.read_csv('climate_change_data.csv')
2 df.head()
```

Out[2]:

	Date	Location	Country	Temperature	CO2 Emissions	Sea Level Rise	Precipitation	Humidity	Wind Speed
0	2000-01-01 00:00:00.000000000	New Williamtown	Latvia	10.688986	403.118903	0.717506	13.835237	23.631256	18.492026
1	2000-01-01 20:09:43.258325832	North Rachel	South Africa	13.814430	396.663499	1.205715	40.974084	43.982946	34.249300
2	2000-01-02 16:19:26.516651665	West Williamland	French Guiana	27.323718	451.553155	-0.160783	42.697931	96.652600	34.124261
3	2000-01-03 12:29:09.774977497	South David	Vietnam	12.309581	422.404983	-0.475931	5.193341	47.467938	8.554563
4	2000-01-04 08:38:53.033303330	New Scottburgh	Moldova	13.210885	410.472999	1.135757	78.695280	61.789672	8.001164

In [3]:

```
1 df['Date'] = pd.to_datetime(df['Date'])
2 df['Month'] = df['Date'].dt.month_name()
```

In [4]:

```
1 df.set_index('Date', inplace=True)
2 df.head()
```

Out[4]:

	Date	Location	Country	Temperature	CO2 Emissions	Sea Level Rise	Precipitation	Humidity	Wind Speed	Month
	2000-01-01 00:00:00.000000000	New Williamtown	Latvia	10.688986	403.118903	0.717506	13.835237	23.631256	18.492026	January
	2000-01-01 20:09:43.258325832	North Rachel	South Africa	13.814430	396.663499	1.205715	40.974084	43.982946	34.249300	January
	2000-01-02 16:19:26.516651665	West Williamland	French Guiana	27.323718	451.553155	-0.160783	42.697931	96.652600	34.124261	January
	2000-01-03 12:29:09.774977497	South David	Vietnam	12.309581	422.404983	-0.475931	5.193341	47.467938	8.554563	January
	2000-01-04 08:38:53.033303330	New Scottburgh	Moldova	13.210885	410.472999	1.135757	78.695280	61.789672	8.001164	January

In [5]:

```
1 df.shape
```

Out[5]: (10000, 9)

In [6]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 10000 entries, 2000-01-01 00:00:00 to 2022-12-31 00:00:00
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Location              10000 non-null  object
1   Country               10000 non-null  object
2   Temperature           10000 non-null  float64
3   CO2 Emissions         10000 non-null  float64
4   Sea Level Rise        10000 non-null  float64
5   Precipitation         10000 non-null  float64
6   Humidity              10000 non-null  float64
7   Wind Speed            10000 non-null  float64
8   Month                 10000 non-null  object
dtypes: float64(6), object(3)
memory usage: 781.2+ KB
```

In [7]:

1 df.describe()

Out[7]:

	Temperature	CO2 Emissions	Sea Level Rise	Precipitation	Humidity	Wind Speed
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	14.936034	400.220469	-0.003152	49.881208	49.771302	25.082066
std	5.030616	49.696933	0.991349	28.862417	28.929320	14.466648
min	-3.803589	182.131220	-4.092155	0.010143	0.018998	0.001732
25%	11.577991	367.109330	-0.673809	24.497516	24.713250	12.539733
50%	14.981136	400.821324	0.002332	49.818967	49.678412	24.910787
75%	18.305826	433.307905	0.675723	74.524991	75.206390	37.670260
max	33.976956	582.899701	4.116559	99.991900	99.959665	49.997664

In [8]:

1 df.describe(include=object)

Out[8]:

	Location	Country	Month
count	10000	10000	10000
unique	7764	243	12
top	North David	Congo	January
freq	12	94	849

In [9]:

1 df.isnull().sum()

Out[9]:

Location	0
Country	0
Temperature	0
CO2 Emissions	0
Sea Level Rise	0
Precipitation	0
Humidity	0
Wind Speed	0
Month	0

dtype: int64

In [10]:

1 df.columns

Out[10]:

```
Index(['Location', 'Country', 'Temperature', 'CO2 Emissions', 'Sea Level Rise',
      'Precipitation', 'Humidity', 'Wind Speed', 'Month'],
      dtype='object')
```

In [11]:

1 df.nunique()

Out[11]:

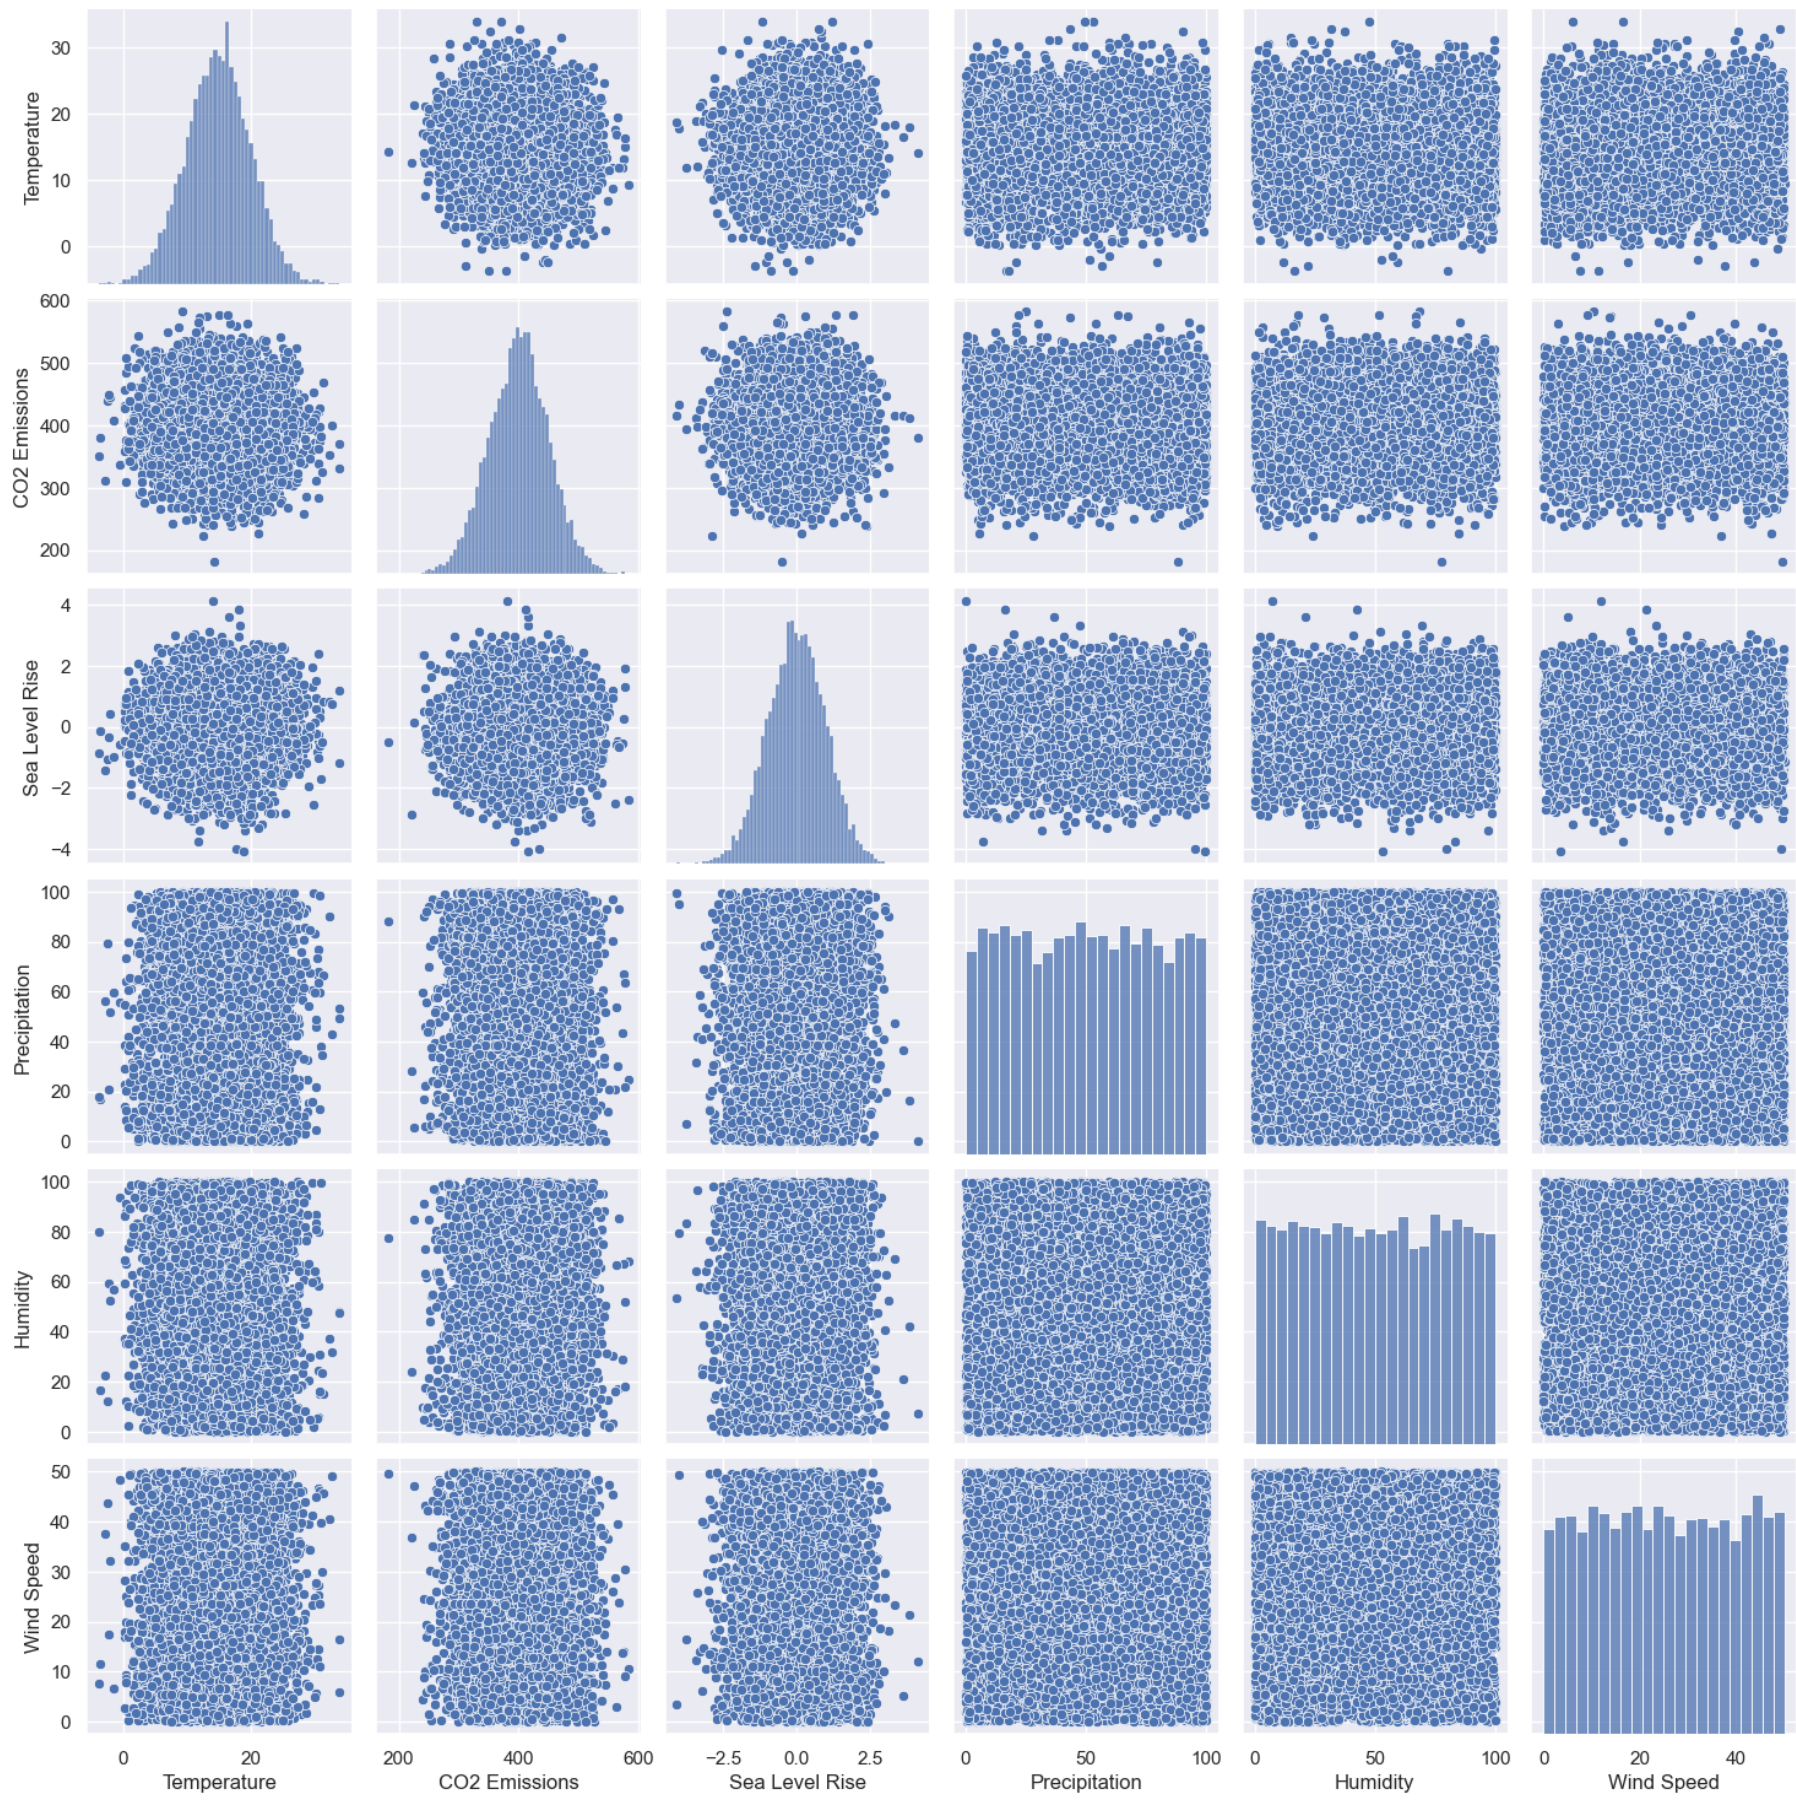
Location	7764
Country	243
Temperature	10000
CO2 Emissions	10000
Sea Level Rise	10000
Precipitation	10000
Humidity	10000
Wind Speed	10000
Month	12

dtype: int64


Data Visualization

```
In [12]: 1 sns.pairplot(df)
```


Out[12]: <seaborn.axisgrid.PairGrid at 0x2670961c880>




```
In [13]: 1 plt.figure(figsize=(10, 6))
2 plt.plot(df['Temperature'], label='Temperature (Celsius)')
3 plt.xlabel('Year')
4 plt.ylabel('Temperature (Celsius)')
5 plt.title('Temperature Trends Over Time')
6 plt.legend()
7 plt.show()
```




```
In [14]: 1 plt.figure(figsize=(10, 6))
2 sns.lineplot(x=df['Month'], y=df['Temperature'])
3 plt.xlabel('Month')
4 plt.ylabel('Temperature (Celsius)')
5 plt.title('Temperature Trends Over Month')
6 plt.show()
```





```
In [15]: 1 plt.figure(figsize=(10, 6))
2 sns.lineplot(x=df.Month, y=df['Humidity'], label='Humidity')
3 sns.lineplot(x=df.Month, y=df['Precipitation'], label='Precipitation')
4
5 plt.xlabel('Month')
6 plt.ylabel('Humidity / Precipitation')
7 plt.title('Humidity & Precipitation Trends Over Month')
8 plt.show()
```




```
In [16]: 1 plt.figure(figsize=(10, 6))
2 sns.lineplot(x=df.Month, y=df['Wind Speed'])
3 plt.xlabel('Month')
4 plt.ylabel('Wind Speed')
5 plt.title('Wind Speed Trends Over Month')
6 plt.show()
```



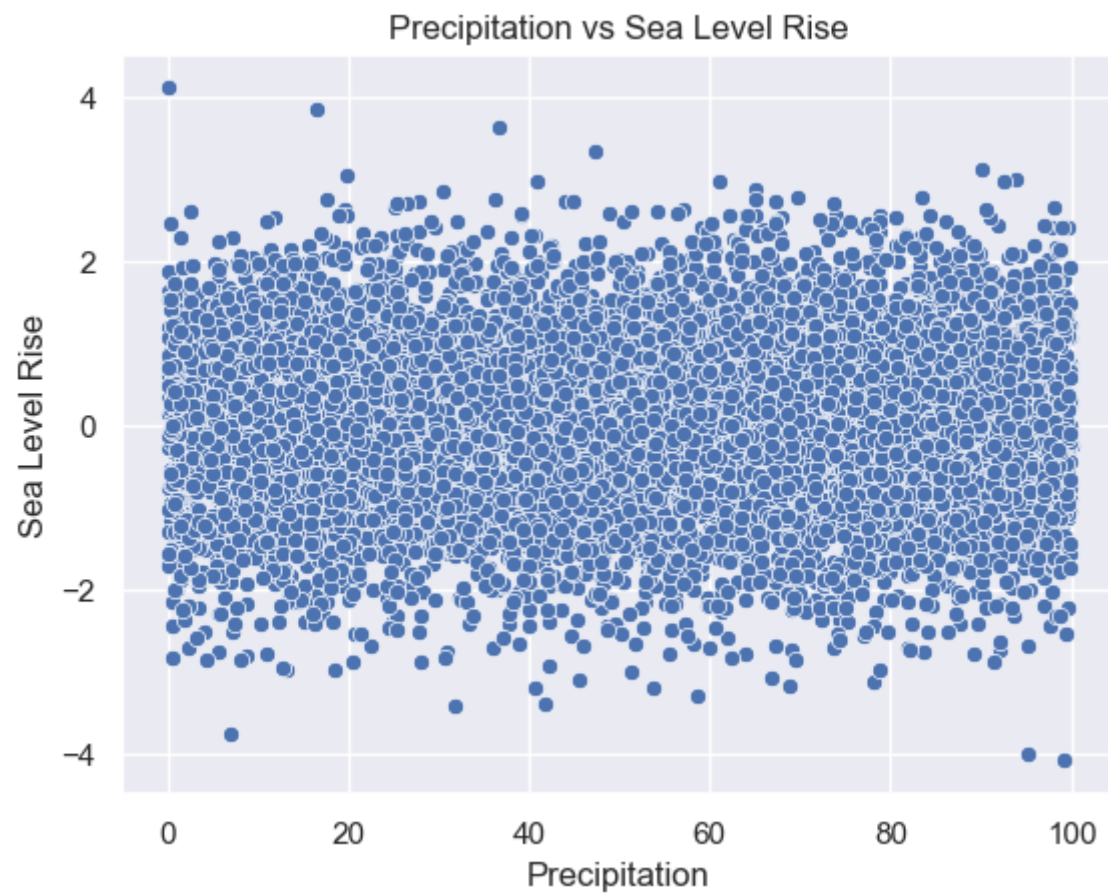
```
In [17]: 1 plt.figure(figsize=(10, 6))
2 sns.lineplot(x=df.Month, y=df['CO2 Emissions'])
3 plt.xlabel('Month')
4 plt.ylabel('CO2 Emissions')
5 plt.title('CO2 Emissions Trends Over Month')
6 plt.show()
```



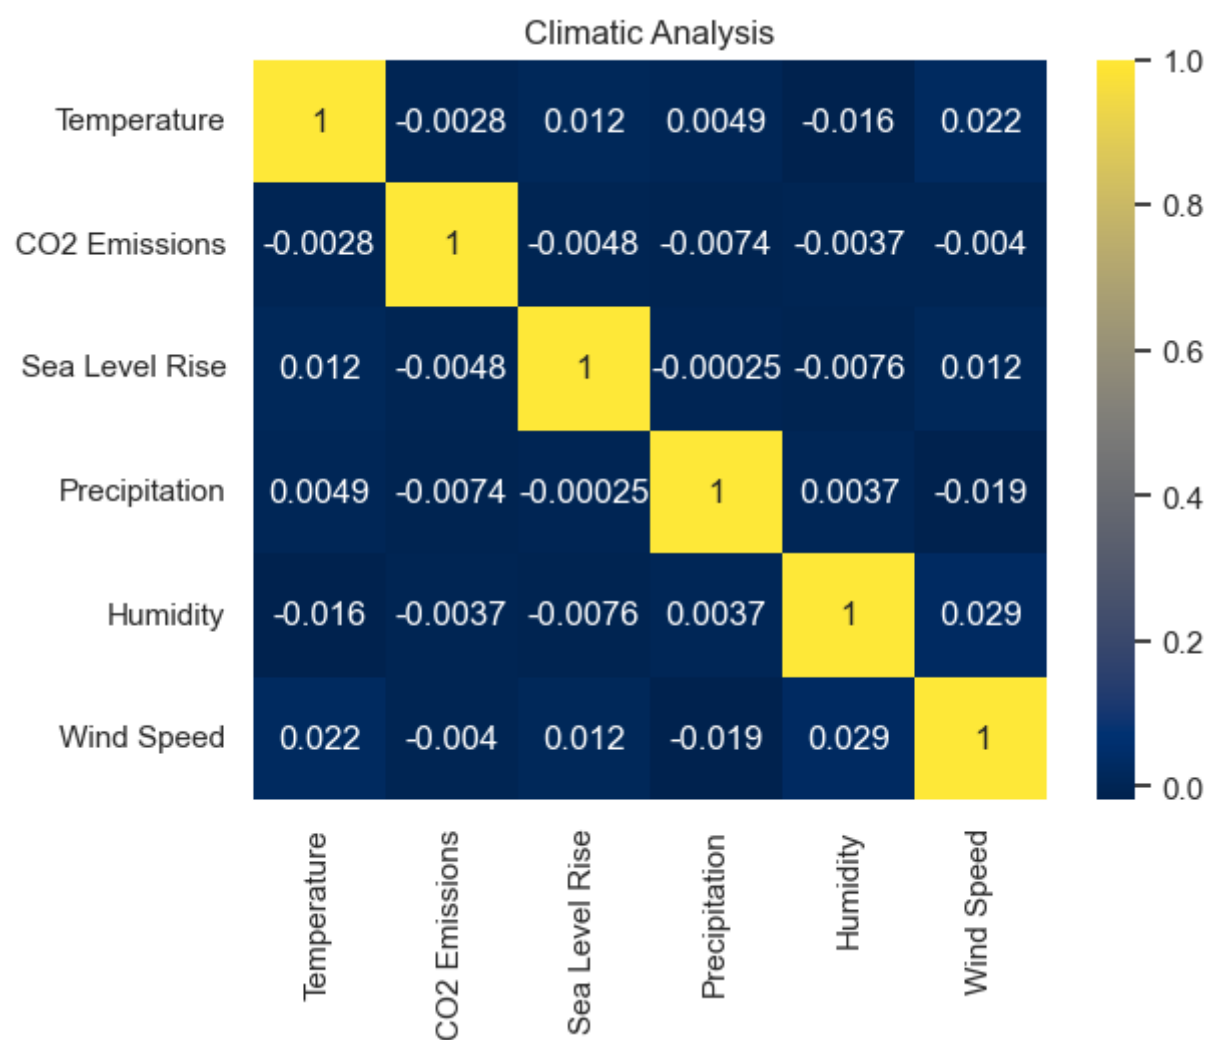
```
In [18]: 1 plt.figure(figsize=(10, 6))
2 sns.lineplot(x=df.index, y=df['Sea Level Rise'])
3 plt.xlabel('Year')
4 plt.ylabel('Sea Level Rise')
5 plt.title('Sea Level Rise Trends Over Time')
6 plt.show()
```



```
In [19]: 1 sns.scatterplot(x='Precipitation', y='Sea Level Rise', data=df)
2 plt.title('Precipitation vs Sea Level Rise')
3 plt.xlabel('Precipitation')
4 plt.ylabel('Sea Level Rise')
5 plt.show()
```



```
In [20]: 1 # Correlation Matrix
2 sns.heatmap(df.corr(), annot=True, cmap='cividis')
3 plt.title('Climatic Analysis')
4 plt.show()
```



Data Preprocessing

```
In [21]: 1 X = df[['Precipitation', 'Humidity', 'Wind Speed']]
2 y = df[['Temperature']]
```

```
In [22]: 1 # Splitting the Dataset
2 from sklearn.model_selection import train_test_split
3 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

Training the Model

- After training the model, the model can be able to predict the temperature according to the given climatic features (i.e., CO2 emissions, sea level rise, precipitation, wind speed, humidity, etc.)

```
In [23]: 1 from sklearn.linear_model import LinearRegression
2 model = LinearRegression()
```

```
In [24]: 1 model.fit(x_train, y_train)
```

```
Out[24]: ▾ LinearRegression
LinearRegression()
```

```
In [25]: 1 len(x_train), len(y_train)
```

```
Out[25]: (8000, 8000)
```

```
In [26]: 1 len(x_test), len(y_test)
```

```
Out[26]: (2000, 2000)
```

```
In [27]: 1 print('Accuracy of the model: ', model.score(x_test, y_test))
```

```
Accuracy of the model:  0.0010176356149836918
```

```
In [28]: 1 from sklearn.metrics import mean_absolute_error, mean_squared_error
2 y_pred = model.predict(x_test)
```

```
In [29]: 1 mae = mean_absolute_error(y_test, y_pred)
2 print(f'Mean Absolute Error: {mae}')
```

```
Mean Absolute Error: 3.999099614445476
```

```
In [30]: 1 mse = mean_squared_error(y_test, y_pred)
2 print(f'Mean Squared Error: {mse}')
```

```
Mean Squared Error: 25.185754765740786
```

```
In [31]: 1 rmse = np.sqrt(mse)
2 print(f'Root Mean Squared Error: {rmse}')
```

```
Root Mean Squared Error: 5.018541099337614
```

```
In [32]: 1 actual_predict= pd.DataFrame({
2     'Actual values': y_test.values.flatten(),
3     'Predicted values': y_pred.flatten()})
4 actual_predict.head()
```

```
Out[32]:
```

	Actual values	Predicted values
0	12.459229	15.095398
1	21.508186	14.960996
2	8.076335	14.936417
3	16.547310	14.933838
4	12.623254	14.809402

```
In [ ]: 1
```